



Network Security

Understanding HTTP/1 and HTTP/2 Method Vulnerabilities

The evolution of HTTP protocols has brought significant performance improvements, but each iteration has also introduced distinct security challenges. This presentation explores the critical vulnerabilities inherent in both HTTP/1.1 and HTTP/2, examining how protocol design decisions create unique attack surfaces that security professionals must understand and mitigate.

HTTP/1: The Text-Based Protocol with Legacy Weaknesses


HTTP/1.1, introduced in 1997, relies on human-readable text-based communication between clients and servers. Whilst this design made debugging simpler during the protocol's early days, it has proven to be a significant security liability in modern web architectures. The plain-text nature of HTTP/1.1 requests and responses creates fundamental parsing ambiguities that attackers have exploited for decades.

The protocol's handling of message boundaries is particularly problematic. HTTP/1.1 uses two distinct mechanisms to determine where one request ends and another begins: the `Content-Length` header, which specifies the exact number of bytes in the message body, and `Transfer-Encoding: chunked`, which allows messages to be sent in smaller pieces without knowing the total size upfront.

Security vulnerabilities emerge when different components in the request chain—load balancers, web application firewalls, reverse proxies, and origin servers—interpret these headers inconsistently. An attacker can craft requests that exploit these interpretation differences, leading to request smuggling attacks where malicious payloads are hidden within legitimate traffic.

☐ Key Vulnerabilities

- Request smuggling via header ambiguity
- HTTP header injection attacks
- Response splitting exploits
- Cache poisoning vectors
- Session hijacking opportunities



HTTP/1 Request Smuggling: A Classic Threat

Request smuggling represents one of the most sophisticated and dangerous vulnerabilities in HTTP/1.1 implementations. This attack vector exploits the fundamental ambiguity that arises when intermediary systems and back-end servers disagree about where one HTTP request ends and the next begins.



CL.TE Attacks

Front-end uses Content-Length whilst back-end uses Transfer-Encoding. Attacker manipulates both headers to create parsing discrepancies, causing the back-end to process attacker-controlled data as the start of the next request.

TE.CL Attacks

Front-end processes Transfer-Encoding whilst back-end relies on Content-Length. The attacker crafts chunked messages that hide malicious requests within seemingly legitimate traffic, bypassing security controls.

3

TE.TE Attacks

Both systems support Transfer-Encoding but can be induced to handle it differently through obfuscation techniques. Subtle variations in header formatting create exploitable parsing differences.

Real-World Impact

- **Account Takeover:** Smuggled requests can hijack user sessions by poisoning the request queue, causing the next user's request to be processed with the attacker's credentials or vice versa.
- **Web Cache Poisoning:** Attackers inject malicious content into cached responses, affecting all subsequent users who retrieve the poisoned cache entry.
- **Security Control Bypass:** Request smuggling can circumvent web application firewalls, authentication mechanisms, and access controls that only inspect the first request in a connection.
- **Credential Harvesting:** Smuggled requests can capture sensitive data from other users' requests, including authentication tokens, session cookies, and personal information.

Detection Challenges

Request smuggling attacks are notoriously difficult to detect because they operate at the protocol level rather than the application level. Traditional security tools often fail to identify these attacks as the malicious payload appears to be part of legitimate traffic.

The attack leaves minimal traces in standard logs, as the smuggled request may be attributed to a different user or session. This makes forensic analysis extremely challenging and allows attackers to maintain persistence whilst evading detection.

HTTP/2: Binary Protocol with New Attack Surfaces

HTTP/2, standardised in 2015, represents a fundamental redesign of HTTP communication. Unlike its text-based predecessor, HTTP/2 employs binary framing, which eliminates many of the parsing ambiguities that plagued HTTP/1.1. The protocol introduces sophisticated features including stream multiplexing, header compression via HPACK, and server push capabilities.

Binary Framing Layer

HTTP/2 breaks down HTTP messages into small, binary-encoded frames. Each frame contains a specific type (DATA, HEADERS, SETTINGS, etc.) and belongs to a particular stream. This structure provides clear message boundaries and eliminates the Content-Length and Transfer-Encoding ambiguities that enabled HTTP/1 smuggling.

Stream Multiplexing

Multiple concurrent request-response exchanges occur over a single TCP connection. Each stream is independent, allowing parallel processing without head-of-line blocking. However, this efficiency creates new attack vectors where malicious actors can exploit stream management mechanisms.

Pseudo-Headers

HTTP/2 introduces special headers prefixed with colons (:method, :path, :scheme, :authority) that replace the HTTP/1.1 request line. These mandatory headers must appear before regular headers, and their improper handling can create security vulnerabilities in implementations.

Design Improvements Over HTTP/1

The binary framing mechanism theoretically prevents traditional request smuggling attacks by providing unambiguous message boundaries. Header compression reduces bandwidth consumption, whilst stream prioritisation allows clients to indicate which resources are most critical. Server push enables proactive content delivery, potentially reducing page load times.

However, these sophisticated features introduce implementation complexity that creates novel vulnerability classes. The protocol's design assumes that all implementations will correctly handle edge cases, but real-world deployments frequently contain subtle flaws that attackers can exploit.

New Attack Considerations

HTTP/2's complexity shifts the security landscape from parsing ambiguities to resource exhaustion and implementation-specific vulnerabilities. The protocol's efficiency features become weapons when abused, allowing small amounts of malicious traffic to have disproportionate impacts on server resources.



HTTP/2 Rapid Reset Attack (CVE-2023-44487)



Attack Statistics

Peak Attack Volume: 398 million requests/second

Amplification Factor: Up to 1000x compared to traditional DDoS

Affected Infrastructure: Google, Amazon, Cloudflare, and major CDN providers

The HTTP/2 Rapid Reset attack, disclosed in October 2023, represents one of the most devastating DDoS techniques ever discovered. This vulnerability exploits a fundamental feature of HTTP/2: the ability to cancel streams using RST_STREAM frames. Whilst stream cancellation serves legitimate purposes—allowing clients to abort unnecessary requests—attackers weaponise this mechanism to create unprecedented server load.

The attack operates by opening numerous streams within a single HTTP/2 connection and immediately cancelling them with RST_STREAM frames. The server expends significant resources processing each stream creation, but receives no corresponding data transmission that would naturally rate-limit the attack. This asymmetry allows attackers to generate millions of requests per second from relatively modest botnets.


Technical Mechanics

- **Stream Creation Flood:** Attackers rapidly open new streams using HEADERS frames, each creating server-side state and consuming processing resources.
- **Immediate Cancellation:** Before the server can respond or process the request fully, the attacker sends RST_STREAM frames to abort each stream.
- **Resource Exhaustion:** The server must allocate memory, update state tables, and potentially execute application logic for each stream, even though they're immediately cancelled.
- **Multiplexing Abuse:** A single TCP connection can carry thousands of these attack streams, making traditional connection-based rate limiting ineffective.

Record-breaking attacks utilising this technique achieved 398 million requests per second—nearly three times larger than any previously recorded DDoS attack. Small botnets with hundreds of nodes proved capable of overwhelming major cloud infrastructure providers, demonstrating the attack's devastating efficiency.

HTTP/2 CONTINUATION Flood Vulnerability (2024)

Discovered in early 2024, the CONTINUATION Flood vulnerability exposes a critical flaw in how HTTP/2 implementations handle header fragmentation. HTTP/2 allows large headers to be split across multiple CONTINUATION frames, with an END_HEADERS flag indicating the final frame. Many implementations failed to properly validate or limit these frame sequences, creating a memory exhaustion attack vector.



Affected Software Ecosystem

- Apache HTTP Server (httpd)**: Versions prior to 2.4.58 vulnerable to unbounded memory allocation during CONTINUATION processing.
- Envoy Proxy**: Critical infrastructure component used by major cloud platforms, affected before emergency patches in March 2024.
- Node.js HTTP/2 Implementation**: Versions before 18.19.1, 20.11.1, and 21.6.2 susceptible to CONTINUATION-based DoS.
- Golang net/http2**: Go's standard library HTTP/2 support required updates to enforce CONTINUATION frame limits.
- Nginx (with nghttp2)**: Deployments using the nghttp2 library for HTTP/2 support needed urgent patching.
- HAProxy**: Load balancer implementations vulnerable when HTTP/2 features enabled.

Impact Assessment

The widespread nature of this vulnerability demonstrated the systemic risks inherent in protocol implementation complexity. Unlike application-specific bugs, this flaw existed in the fundamental HTTP/2 processing code used across the entire web infrastructure stack.

Proof-of-concept exploits demonstrated that single attackers could crash major web servers with minimal traffic, as each connection could spawn multiple malicious streams. The lack of default frame limits in many implementations meant that servers would faithfully buffer megabytes or gigabytes of header data before recognising an attack.



HTTP/2 Desynchronisation & Request Smuggling via Downgrading

Whilst HTTP/2's binary framing eliminates traditional request smuggling within HTTP/2-only infrastructure, real-world deployments frequently involve protocol downgrading. Many architectures use HTTP/2 for client-facing connections to leverage performance benefits, but communicate with back-end services via HTTP/1.1 for compatibility reasons. This architectural pattern reintroduces smuggling vulnerabilities with novel characteristics.

01

HTTP/2 Binary Reception

Front-end proxy receives HTTP/2 binary frames from client, parsing headers and pseudo-headers according to HTTP/2 specifications including stream multiplexing and HPACK compression.

02

Protocol Translation

Proxy translates HTTP/2 frames into HTTP/1.1 text format for back-end communication. This conversion process makes critical decisions about header ordering, special character handling, and message boundary encoding.

03

Parsing Inconsistency Exploitation

Attacker crafts HTTP/2 requests that translate into ambiguous HTTP/1.1 representations. The front-end and back-end systems interpret the resulting HTTP/1.1 messages differently, creating desynchronisation.

04

Malicious Request Injection

Smuggled request content executes with the privileges of the next legitimate user, enabling session hijacking, cache poisoning, or unauthorised access to sensitive resources.

Attack Vectors

- **Header Name Manipulation:** HTTP/2 allows broader character sets in header names. Translators may handle special characters inconsistently, creating smuggling opportunities.
- **Pseudo-Header Abuse:** Attackers inject multiple :method or :path pseudo-headers, exploiting variations in how translators prioritise or merge them.
- **Content-Length Injection:** Crafting HTTP/2 requests that generate conflicting Content-Length values during translation to HTTP/1.1.
- **Transfer-Encoding Smuggling:** HTTP/2 prohibits Transfer-Encoding, but attackers may inject it as a regular header, causing confusion during downgrade.


Real-World Exploitation

Security researchers have demonstrated practical attacks against major CDN providers and cloud platforms. In one notable case, attackers achieved request smuggling through protocol downgrade, successfully hijacking user sessions and accessing administrative interfaces.

The complexity of detecting these attacks stems from the fact that the HTTP/2 traffic appears entirely legitimate when inspected at the front-end. Only by analysing the resulting HTTP/1.1 traffic—which often occurs in internal, less-monitored network segments—can defenders identify the smuggled requests.



Comparative Summary: HTTP/1 vs HTTP/2 Vulnerabilities



HTTP/1.1 Vulnerability Profile

Primary Weakness: Text-based parsing ambiguities create fundamental security challenges that are difficult to eliminate without protocol redesign.

- Request smuggling via Content-Length/Transfer-Encoding conflicts
- HTTP header injection enabling response splitting
- Cache poisoning through ambiguous message boundaries
- Session hijacking via request queue manipulation

Root Cause: Multiple valid interpretations of the same byte stream allow attackers to exploit parsing differences between security devices and application servers.



Http/2

HTTP/2 Vulnerability Profile

Primary Weakness: Protocol complexity and resource management mechanisms create novel attack surfaces focused on availability and resource exhaustion.

- Rapid Reset attacks enabling record-breaking DDoS
- CONTINUATION Flood causing memory exhaustion
- Desynchronisation attacks during HTTP/2-to-HTTP/1 downgrading
- Stream multiplexing abuse overwhelming server resources

Root Cause: Performance optimisations and sophisticated features introduce implementation complexity that creates exploitable edge cases and resource exhaustion vectors.

Attack Sophistication

HTTP/1: Exploits require understanding of text parsing but are conceptually straightforward.

HTTP/2: Attacks leverage complex protocol features, requiring deeper technical knowledge but offering greater impact.

Detection Difficulty

HTTP/1: Smuggling attacks hide within legitimate traffic, challenging to detect without deep packet inspection.

HTTP/2: Resource exhaustion creates observable impact, but attribution and prevention remain difficult.

Mitigation Complexity

HTTP/1: Requires strict parsing enforcement and consistent header handling across infrastructure.

HTTP/2: Demands regular patching, resource limits, and careful protocol downgrade management.

Both protocols require vigilant security posture, continuous monitoring, and proactive patching. The choice between HTTP/1.1 and HTTP/2 should weigh performance benefits against the security team's capacity to address protocol-specific vulnerabilities.

Mitigation Strategies & Best Practices

Defending against HTTP protocol vulnerabilities requires a multi-layered approach combining configuration hardening, infrastructure updates, monitoring enhancements, and architectural considerations. Security teams must address both HTTP/1 legacy risks and HTTP/2 novel attack vectors simultaneously.

Protocol Configuration & Hardening

- **Disable HTTP/2 Selectively:** If infrastructure contains unpatched HTTP/2 implementations, temporarily disable the protocol until updates can be deployed. Assess performance impact before making permanent decisions.
- **Enforce Strict Parsing:** Configure web servers and proxies to reject ambiguous requests. Enable strict header validation, disallow conflicting Content-Length/Transfer-Encoding headers, and reject malformed pseudo-headers.
- **Limit Frame and Header Sizes:** Set conservative limits on HTTP/2 CONTINUATION frames, maximum header list size, and concurrent stream counts to prevent resource exhaustion attacks.
- **Disable Protocol Downgrading:** When possible, maintain HTTP/2 end-to-end throughout the infrastructure stack. If downgrading is necessary, deploy specialised reverse proxies with proven security records in protocol translation.

Infrastructure & Architecture

- **Deploy DDoS Protection Services:** Utilise cloud-based DDoS mitigation platforms (Cloudflare, AWS Shield, Akamai) that can absorb Rapid Reset floods before they reach origin infrastructure. These services employ specialised detection and filtering for HTTP/2 attacks.
- **Implement Request Rate Limiting:** Apply rate limits at multiple levels: per-IP, per-connection, and per-stream. HTTP/2 requires stream-level rate limiting, not just connection-level controls.
- **Separate Critical Services:** Isolate high-value applications on dedicated infrastructure with enhanced security controls, preventing attacks on public-facing services from impacting critical systems.
- **Deploy Web Application Firewalls:** Modern WAFs with HTTP/2 awareness can detect anomalous frame patterns, excessive stream cancellations, and suspicious header sequences.

Software Updates & Patch Management

- **Prioritise HTTP/2 Security Updates:** Treat HTTP/2-related vulnerabilities as critical. Patches for Rapid Reset, CONTINUATION Flood, and similar issues require immediate deployment.
- **Maintain Current Software Versions:** Ensure all components in the HTTP processing chain run patched versions: Apache HTTP Server ($\geq 2.4.58$), Nginx with `nghttp2` updates, Node.js ($\geq 18.19.1, \geq 20.11.1, \geq 21.6.2$), Golang (latest stable), Envoy Proxy (current release).
- **Automate Security Updates:** Implement automated patch management for non-critical systems, with rapid deployment procedures for emergency security fixes.
- **Test Updates in Staging:** Validate patches in non-production environments before deploying to live systems, but expedite this process for critical security fixes.

Monitoring & Detection

- **Monitor HTTP/2 Frame Patterns:** Establish baselines for normal RST_STREAM, CONTINUATION, and other frame types. Alert on deviations indicating attacks.
- **Track Resource Utilisation:** Monitor server memory, CPU, and connection counts. Rapid increases may indicate ongoing attacks.
- **Analyse Traffic Anomalies:** Implement machine learning-based anomaly detection to identify novel attack patterns that signature-based systems might miss.
- **Enable Comprehensive Logging:** Log HTTP/2 frame types, stream counts, and protocol downgrade events. Ensure logs capture sufficient detail for forensic analysis without overwhelming storage.


Conclusion: Navigating the Evolving HTTP Security Landscape

The evolution from HTTP/1.1 to HTTP/2 demonstrates a fundamental truth about protocol design: optimising for performance and efficiency often introduces new security challenges. HTTP/2's binary framing successfully eliminated text-based parsing ambiguities, but the protocol's sophisticated features—stream multiplexing, header compression, and server push—created novel attack surfaces that continue to challenge defenders.

Understanding the distinct vulnerability profiles of each protocol is essential for effective defence. HTTP/1.1's text-based design enables request smuggling through parsing ambiguities, whilst HTTP/2's resource management mechanisms facilitate devastating DDoS attacks like Rapid Reset and CONTINUATION Flood. Organisations employing mixed protocol environments face the additional challenge of securing protocol downgrade scenarios, where HTTP/2's benefits can reintroduce HTTP/1's classic vulnerabilities.

Key Takeaways for Security Teams

- Protocol Awareness is Critical:** Security tools and personnel must understand protocol-specific attack vectors. Traditional security controls designed for HTTP/1 may prove ineffective against HTTP/2 exploits.
- Architecture Matters:** Protocol downgrade scenarios create hybrid vulnerabilities. End-to-end HTTP/2 or carefully secured translation layers are preferable to ad-hoc mixed deployments.
- Proactive Defence is Essential:** Waiting for attacks to occur is insufficient. Regular patching, configuration hardening, and monitoring must be continuous processes.



398M	1000X
Requests/Second	Amplification
Peak HTTP/2 Rapid Reset attack volume recorded	Attack efficiency compared to traditional DDoS methods
2024	
CONTINUATION	
Year critical HTTP/2 memory exhaustion vulnerability disclosed	

Looking Ahead: HTTP/3 and Beyond

The future of web communication lies in HTTP/3, which replaces TCP with QUIC (Quick UDP Internet Connections). This next-generation protocol promises improved performance, enhanced security through mandatory encryption, and better resilience to packet loss. However, the lessons from HTTP/2's security journey must inform HTTP/3's deployment: sophisticated features require equally sophisticated security analysis, and new protocols should be adopted cautiously with full awareness of potential vulnerabilities.

"Security is not a product, but a process. As HTTP protocols evolve to meet performance demands, our defensive strategies must evolve in parallel. Understanding protocol internals, maintaining vigilant patch management, and implementing defence-in-depth remain the cornerstone of web security."

The web's security posture depends on collective action: software vendors must prioritise security in protocol implementations, infrastructure operators must deploy updates promptly, and security researchers must continue identifying and disclosing vulnerabilities responsibly. Only through this collaborative approach can we build a faster, more efficient, and more secure web for all users.